

Machine Problem 2

Handed Out: Oct 12nd, 2018

Due: Nov 5, 2018 (11:59pm)

TA: Ashutosh Dhekne

Abstract

This machine problem tests your understanding of reliable packet transfer. You will use UDP to implement your own version of TCP. Your implementation must be able to tolerate packet drops, allow other concurrent connections a fair chance, and must not be overly nice to other connections (should not give up the entire bandwidth to other connections).

1 Introduction

In this MP, you will implement a transport protocol with properties equivalent to TCP. You have been provided with a file called `sender_main.c`, which declares the function `void reliablyTransfer(char* hostname, unsigned short int hostUDPPort, char* filename, unsigned long long int bytesToTransfer)`. This function should transfer the first `bytesToTransfer` bytes of `filename` to the receiver at `hostname:hostUDPPort` correctly and efficiently, even if the network drops or reorders some of your packets. You also have `receiver_main.c`, which declares `void reliablyReceive(unsigned short int myUDPPort, char* destinationFile)`. This function is `reliablyTransfer`'s counterpart, and should write what it receives to a file called `destinationFile`.

2 What is expected in this MP?

Your job is to implement `reliablyTransfer()` and `reliablyReceive()` functions, with the following requirements:

- The data written to disk by the receiver must be exactly what the sender was given.
- Two instances of your protocol competing with each other must converge to roughly fairly sharing the link (same throughputs $\pm 10\%$), within 100 RTTs. The two instances might not be started at the exact same time.
- Your protocol must be somewhat TCP friendly: an instance of TCP competing with you must get on average at least half as much throughput as your flow.
- An instance of your protocol competing with TCP must get on average at least half as much throughput as the TCP flow. (Your protocol must not be overly nice.)

- All of the above should hold in the presence of any amount of dropped packets. All flows, including the TCP flows, will see the same rate of drops. The network will not introduce bit errors.
- Your protocol must, in steady state (averaged over 10 seconds), utilize at least 70% of bandwidth when there is no competing traffic, and packets are not artificially dropped or re-ordered.
- You cannot use TCP in any way. Use SOCK_DGRAM (UDP), not SOCK_STREAM.

The test environment has a 40Mbps connection, and a 20ms RTT.

3 VM Setup - Replicating the Test Environment

You'll need 2 VMs to test your client and server together. Unfortunately, VirtualBox's default setup does not allow its VMs to talk to the host or each other. There is a simple fix, but then that prevents them from talking to the internet. So, be sure you have done all of your apt-get installs before doing the following! (To be sure, just run: `sudo apt-get install gcc make gdb valgrind iperf tcpdump`) Make sure the VMs are fully shut down. Go to each of their Settings menus, and go to the Network section. Switch the Adapter Type from NAT to "host-only", and click ok. When you start them, you should be able to ssh to them from the host, and it should be able to ping the other VM. You can use `ifconfig` to find out the VMs' IP addresses. If they both get the same address, `sudo ifconfig eth0 newipaddr` will change it. (If you make the 2 nd VM by cloning the first + choosing reinitialize MAC address, that should give different addresses.)

New in MP2: You can use the same basic test environment described above. However, the network performance will be ridiculously good (same goes for testing on localhost), so you'll need to limit it. The autograder uses `tc` . If your network interface inside the VM is `eth0`, then run (from inside the VM) the following command:

```
sudo tc qdisc del dev eth0 root 2>/dev/null
```

to delete existing tc rules. Then use,

```
sudo tc qdisc add dev eth0 root handle 1:0 netem delay 20ms loss 5%
```

followed by

```
sudo tc qdisc add dev eth0 parent 1:1 handle 10: tbf rate 40Mbit burst 10mb latency 1ms
```

will give you a 40Mbit, 20ms RTT link where every packet sent has a 5% chance to get dropped. Simply omit the `loss n%` part to get a channel without artificial drops.

(You can run these commands just on the sender; running them on the receiver as well won't make much of a difference, although you'll get a 40ms RTT if you don't adjust the delay to account for the fact that it gets applied twice.)

4 Autograder and Submission

Similar to the MP1, checkout your GIT directory from the class release repository. The contents of the checked out mp2 folder will be **different** from mp1 when you first checked it out. Use these new programs as a starting point and make the modifications required for this assignment.

The following lists the **additional notes for MP2**:

- Update the partners.txt file with your netIDs. Only update for one of the partners.
- The MTU on the test network is 1500, so up to a 1472 byte payload (IPv4 header is 20 bytes, UDP is 8) won't get fragmented. You can `sendto()` larger packets and the sockets library's UDP will handle fragmentation/reassembly for you. It's up to you to reason out the benefits and drawbacks of using large UDP packets in various settings.
- You can use the provided main files to be sure your program runs with the right interface, or write your own. **Executable names:** `reliable_sender` and `reliable_receiver`. As with mp1, a single run of "make" (with no arguments) inside your mp2 directory should build both binaries.
- Be sure you have a clean design for implementing the send/receive buffers. Trying to figure out where to get the data to resend an old packet won't be fun if your send window's buffer doesn't have a nice clean interface.
- Input files on the grader are READ-ONLY. Do not use the "rb+" mode to read them; the '+' adds write permission. (In general, you shouldn't use "rb+" unless you need it).
- Input files on the grader are general binary data, NOT text.

Modify the Makefile such that a simple `make` command in your mp2 folder creates the required executables. The autograder does just that. Be careful about the executable filenames and output filenames.

Submission instructions are same as for mp1. Tests generally take **10-15 minutes**, and there may be a queue of students. Expect 2 hour+ waits near the deadline (but don't worry, you will be graded based on when you entered the queue).

PLEASE do not fall into the trap of "debugging on the autograder". If you submit a new version every time you make some change that might help pass an extra test, you are going to waste a lot of time waiting for results. Rather, only submit when you have made major progress or have definitively figured out what you were previously doing wrong. If you aren't genuinely surprised that your most recent submission didn't increase your score, you are submitting too often. Your grade is the highest score that the auto-grader ever gives you.

Do refer to MP1 instructions for other notes.